

# Musical Instrument Classification Using a Hybrid Neural Network

**Landon Buell**

B.S. Physics, Dec 2020

M.S. Computer Science, Exp. 2024

University of New Hampshire

**Kevin Short**

University Professor & Professor of Mathematics

Founder, Integrated Applied Mathematics Program

University of New Hampshire

8 Jan 2021

# Presentation Outline

**Introduce** the *problem* that we are going to solve



**Develop** *Neural Networks* as the solution



**Discuss** consequences and *improvements* to the solution



**Analyze** the *performance* of the improvements

# Introduction

# Mapping Sounds to Sources



Humans are  
proficient at mapping  
sounds to sources



Impractical at a large  
scale



Computers are not  
proficient at mapping  
sounds to sources



Can handle large  
volumes of data

“Birds inspired us to fly, burdock plants inspired Velcro and nature has inspired many other inventions. It seems only logical then, to look to the brain’s architecture for inspiration on how to build an intelligent machine.”

- Aurelion Geron, Former YouTube Video Classification lead

# The Neural Network

# Structure

Consider a neural network to be just like a mathematical function

Composed of smaller functions called *layers*

Transform *features* into *predictions*



*Inputs* are properties of digital audio files from London's *Philharmonia* Orchestra and University of Iowa's *Electronic Music Studios*



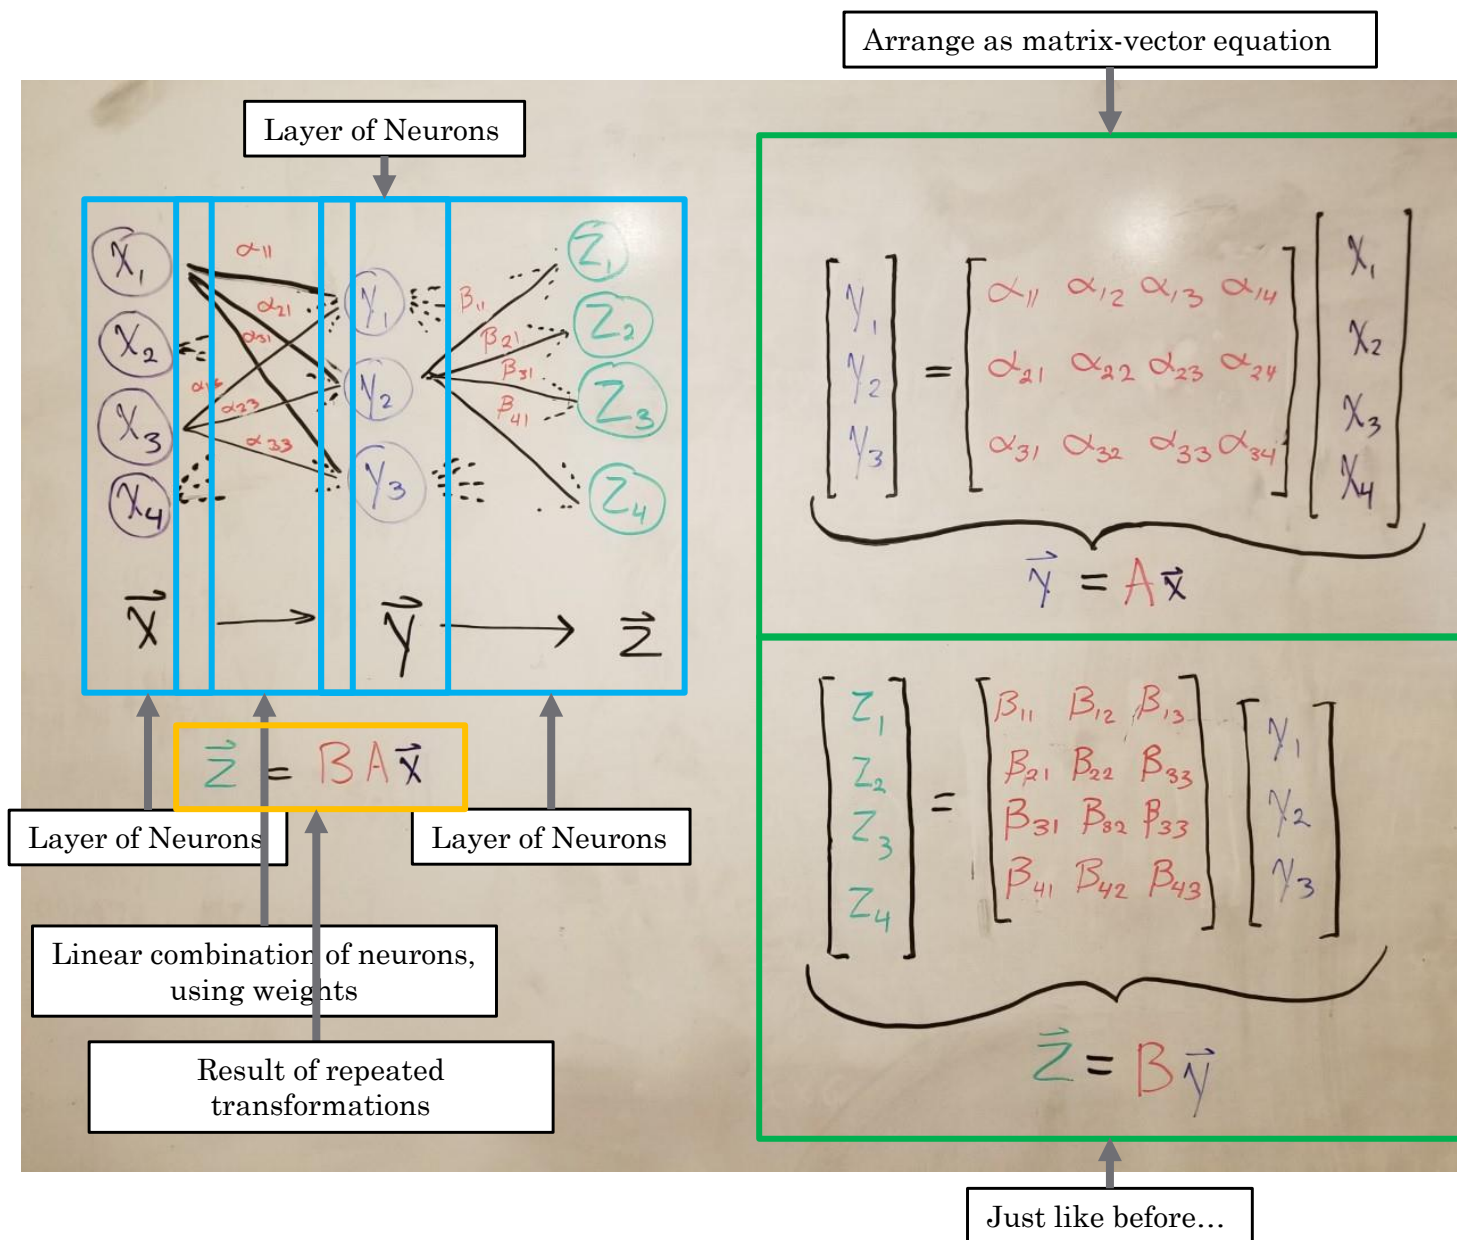
*Outputs* are integers that correspond to musical instruments



We group samples with similar input properties

## Classification Model



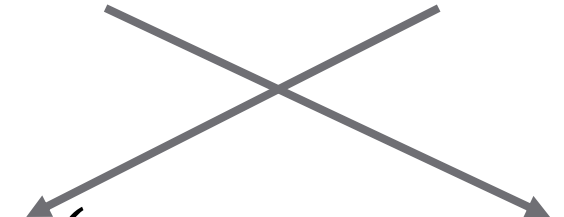


# The Multilayer Perceptron (MLP)

- Connect layers with *weights*
- *Dense* Connectivity
- Handles 1D samples

# MLP (Cont.)

- Formal implementations add a *bias* and an *activation function*

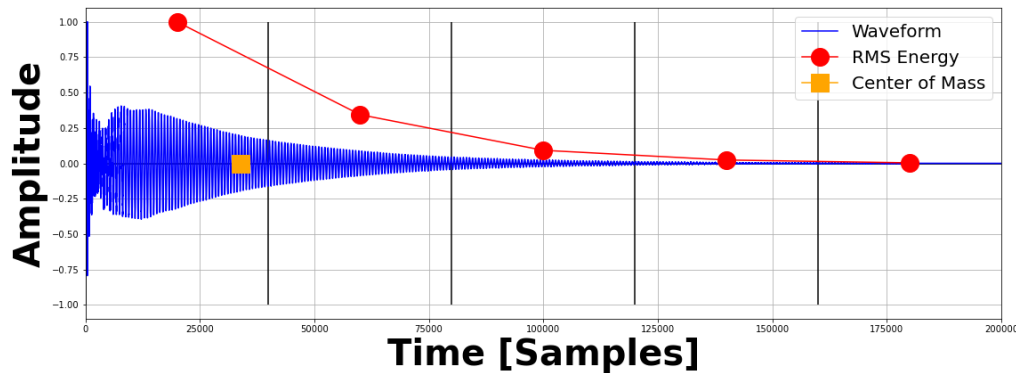
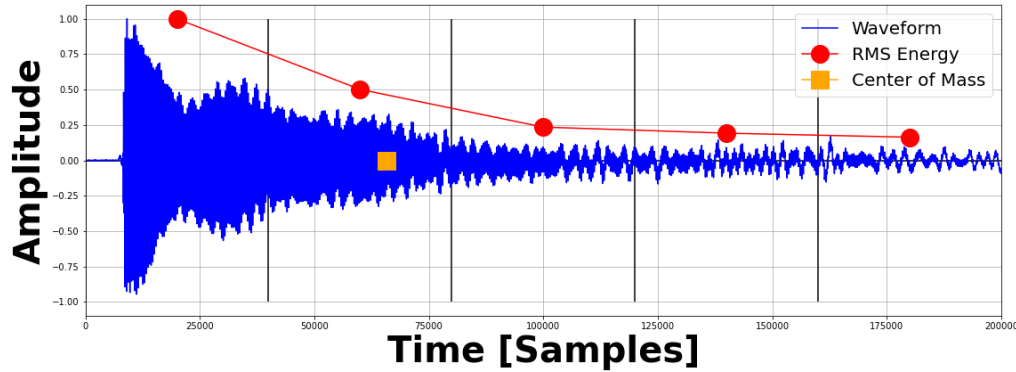


The diagram consists of two gray arrows originating from the text 'bias' and 'activation function' in the list item above. One arrow points to the  $\vec{b}^{(l)}$  term in the equation, and the other points to the  $\sigma^{(l)}$  term.

$$\vec{x}^{(l)} = \sigma^{(l)} \left( \underbrace{W^{(l)} \vec{x}^{(l-1)}}_{\text{Matrix-vector product}} + \vec{b}^{(l)} \right)$$

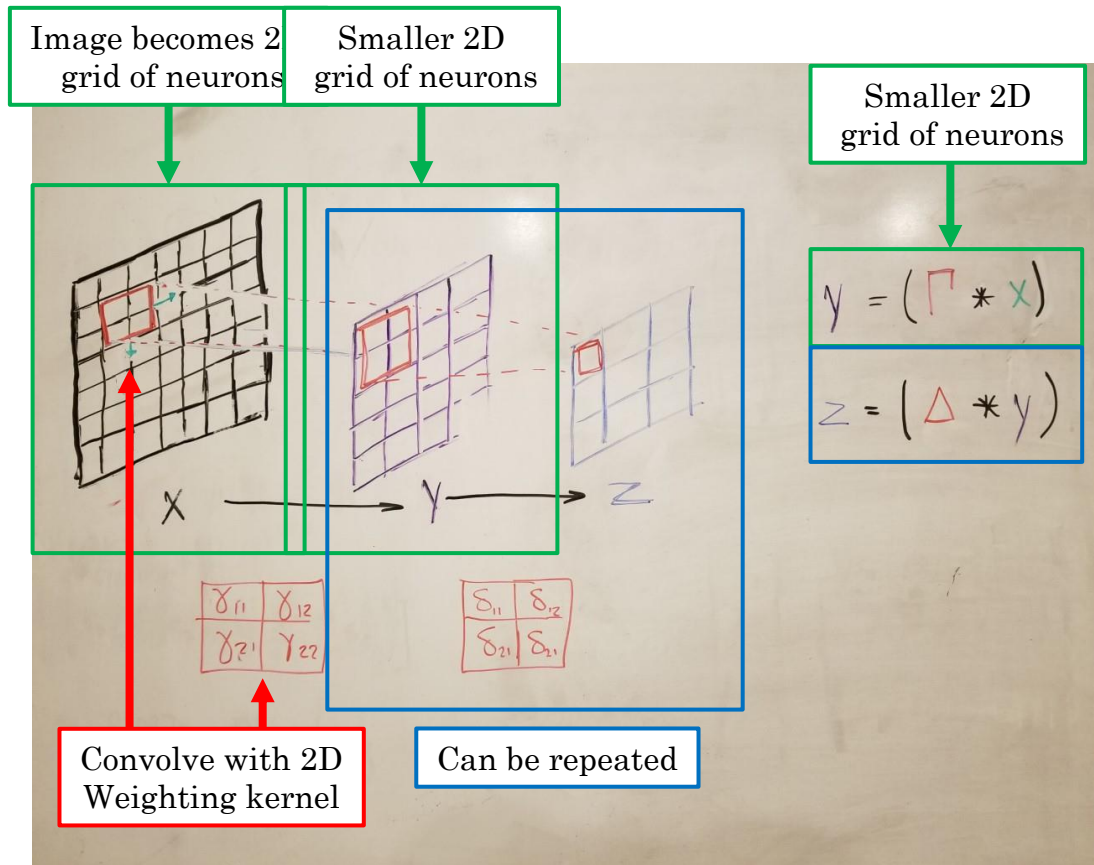
- Allows us to model more complex decision boundaries in real-world problems

# Features for the MLP



- Time Domain Envelope (x5)
- Zero Crossing Rate
- Temporal Center of Mass
- Auto Correlation Coefficients (x4)
- Mel Frequency Cepstrum Coefficients (x12)
- Frequency Center of Mass

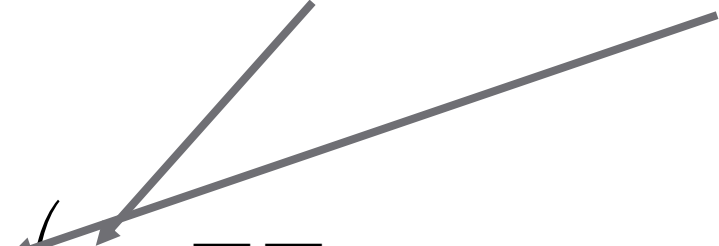
# The Convolutional Neural Network (CNN)



- Connect layers with weighting kernels
- *Sparse* connectivity
- Handles 2D “grid” of neurons (images)
- Followed by layer of pooling

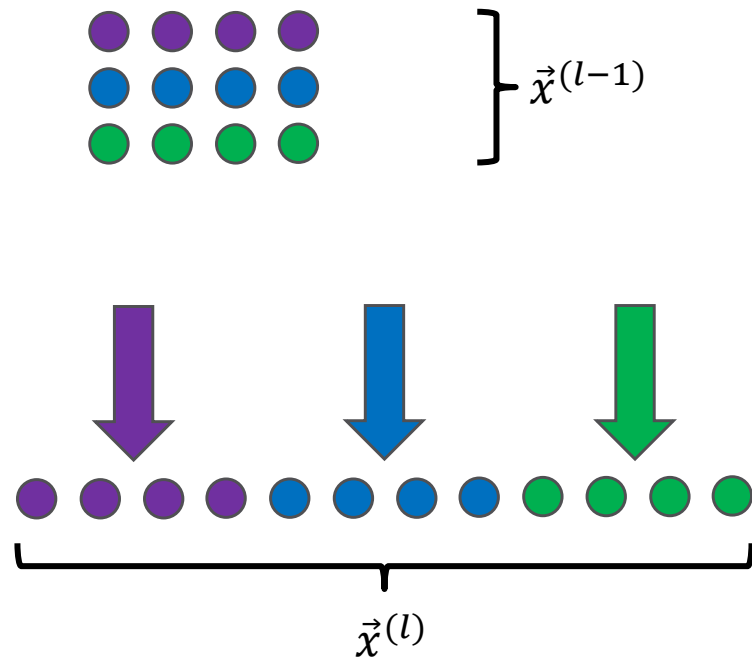
# CNN (Cont.)

- Formal implementation also add a *bias* and an *activation function*

$$x^{(l)}[i,j] = \sigma^{(l)} \left( b^{(l)} + \underbrace{\sum_u \sum_v W^{(l)}[i,j] x^{(l-1)}[i-u, j-v]}_{\text{2D Convolution operation}} \right)$$


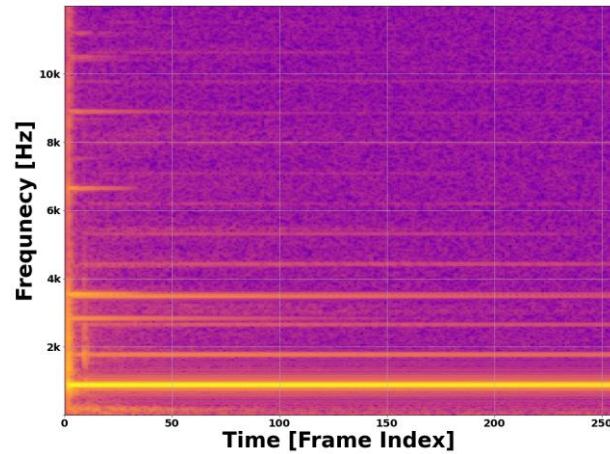
- Detect patterns and features in images, using compressed versions of a parent image

# The Flattening Operation

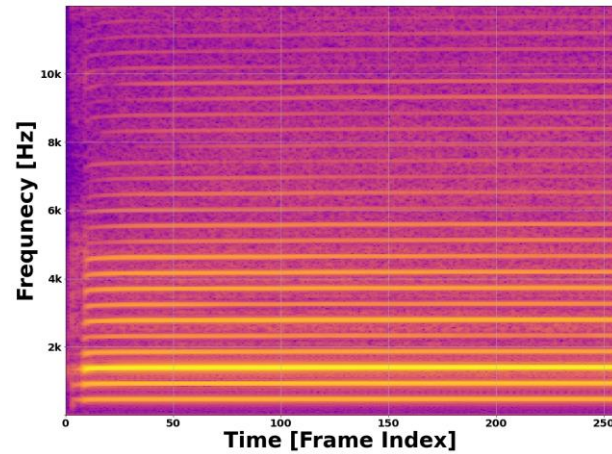


- Transforms 2D grid of neurons to a 1D row/col of neurons
- Allows images to be passed into dense layers

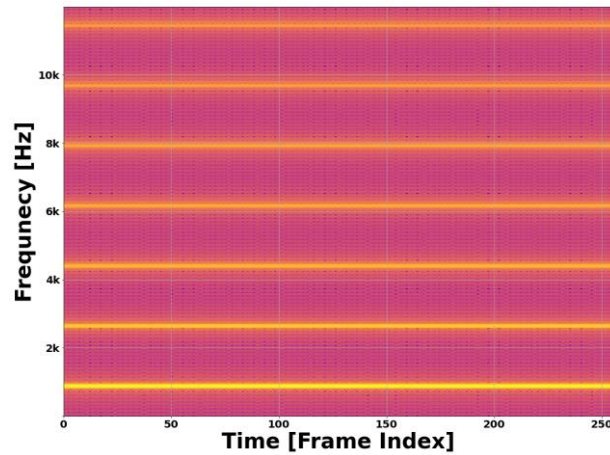
Vibraphone - A5



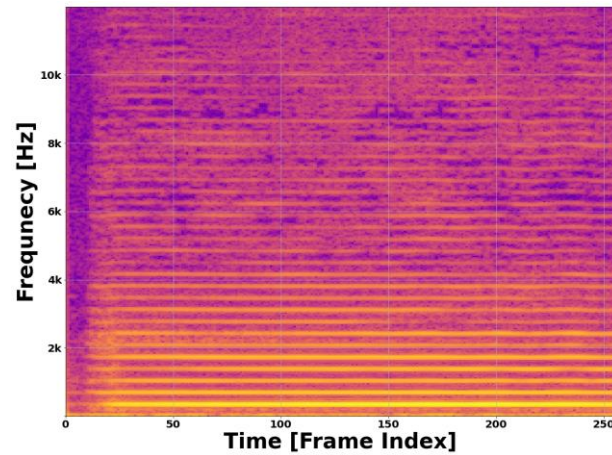
Oboe – Bb4



Square Wave - A5



Cello – F4



Features  
for the  
CNN

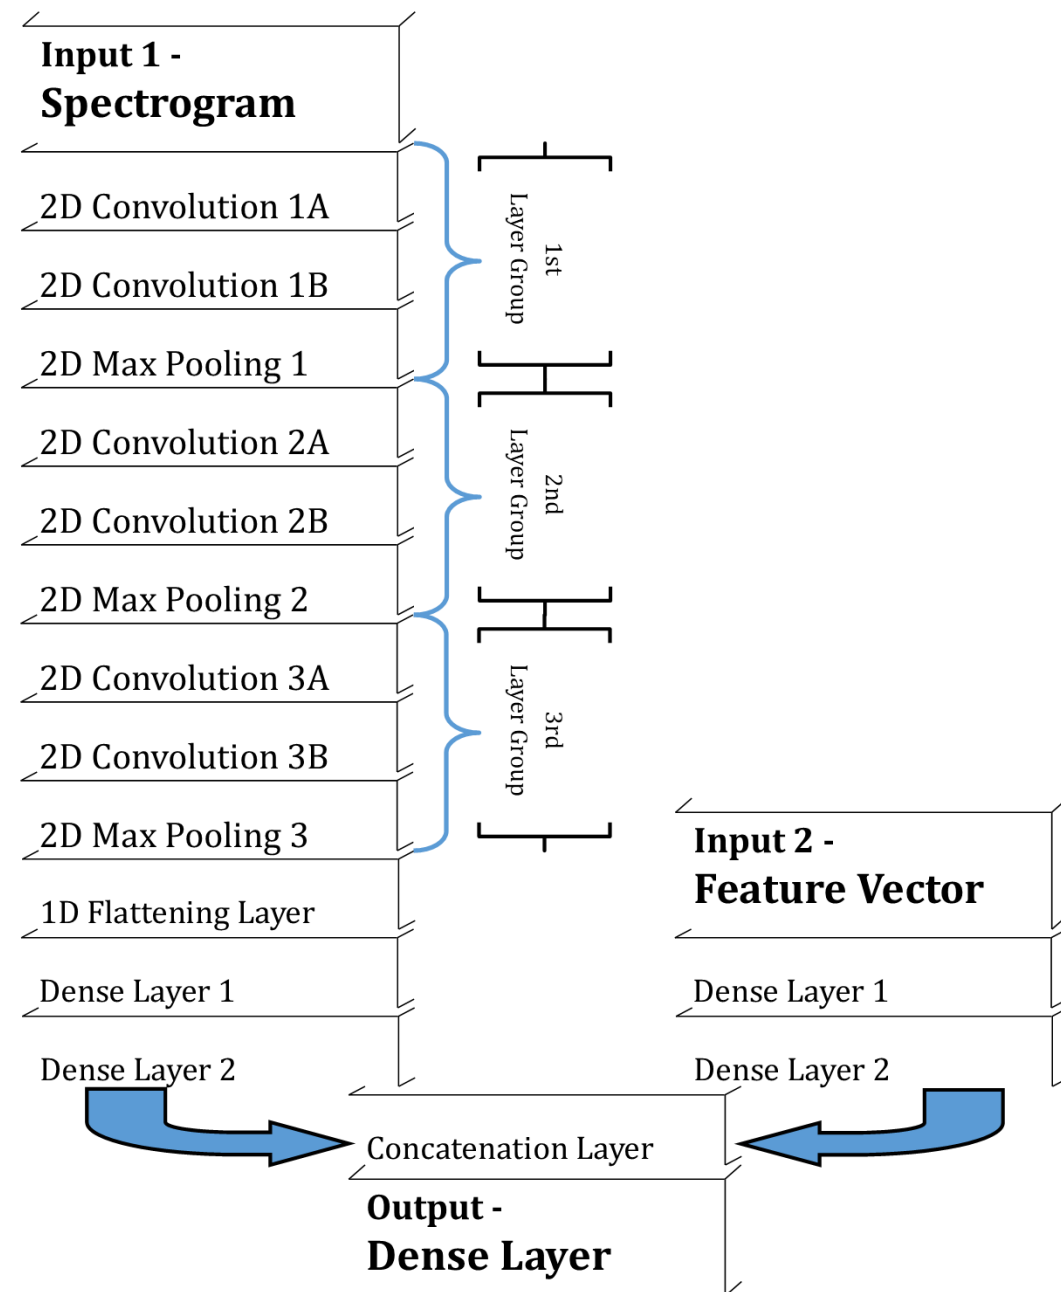
- Both architectures have experimentally shown individual success
- MLP and CNN use different forms of inputs that are initially *incompatible*
  - 1D vs. 2D Inputs
- We can *combine* them to form a *Hybrid* Neural Network
  - Connections made at hidden layer

# Consequence of the Solutions



# Hybrid Network Architecture

# Hybrid Neural Network Architecture



Create each input Branch as its own model

Create combined output layer, wrap as single model instance

Return the hybrid network instance

```
200
201 @staticmethod
202 def MultiInputNetwork (name,inputA,inputB,n_classes):
203     """
204     Create Multi-Input layer neural Network
205     -----
206     name (str) : Name to use for neural network
207     inputA (iter) : list-like of ints indicating input shape of CNN branch
208     inputA (iter) : list-like of ints indicating input shape of MLP branch
209     n_classes (int) : Number of unique output classes
210     -----
211     Return complied tf.keras model
212     """
213     modelCNN = NeuralNetworkModels.ConvolutionalNeuralNetwork2D(inputA,n_classes,
214         filterSizes=[32,32,32],kernelSizes=[(3,3),(3,3),(3,3)],
215         poolSizes=[(3,3),(3,3),(3,3)],neurons=[64,64])
216     modelMLP = NeuralNetworkModels.MultilayerPerceptron(inputB,n_classes)
217
218     x = keras.layers.concatenate([modelCNN.output,modelMLP.output])
219
220     x = keras.layers.Dense(units=n_classes,activation='softmax',name='Output')(x)
221     modelMain = keras.Model(name=name,inputs=[modelCNN.input,modelMLP.input],outputs=x)
222
223     modelMain.compile(optimizer=keras.optimizers.Adam(learning_rate=0.001,
224         beta_1=0.9,beta_2=0.999,epsilon=1e-07),
225         loss=keras.losses.CategoricalCrossentropy(),
226         metrics=[keras.metrics.Accuracy(),keras.metrics.Precision(),keras.metrics.Recall()])
227
228     return modelMain
```

Concatenation of each model's output activations

Compile the model, prepare for training/testing

# Implementation (Tensorflow.keras)



This is called *Multimodal*  
or *Multiview* Deep  
Learning



Transform incompatible  
inputs to a compatible  
format at an internal layer

# Performance of the Hybrid Model

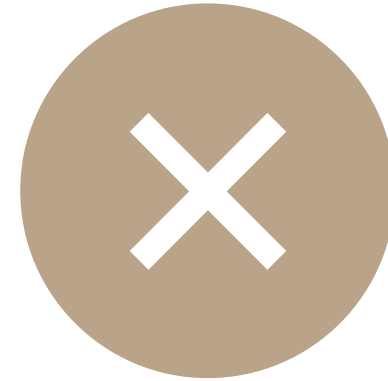
# Evaluating a Model



WE *TRAIN* A MODEL ON A  
*SUBSET* OF ALL SAMPLES



THE MODEL *LEARNS* A SET OF  
PARAMETERS IN EACH LAYER  
THAT ALLOWS THE MAPPING OF  
*FEATURES* TO *PREDICTIONS*

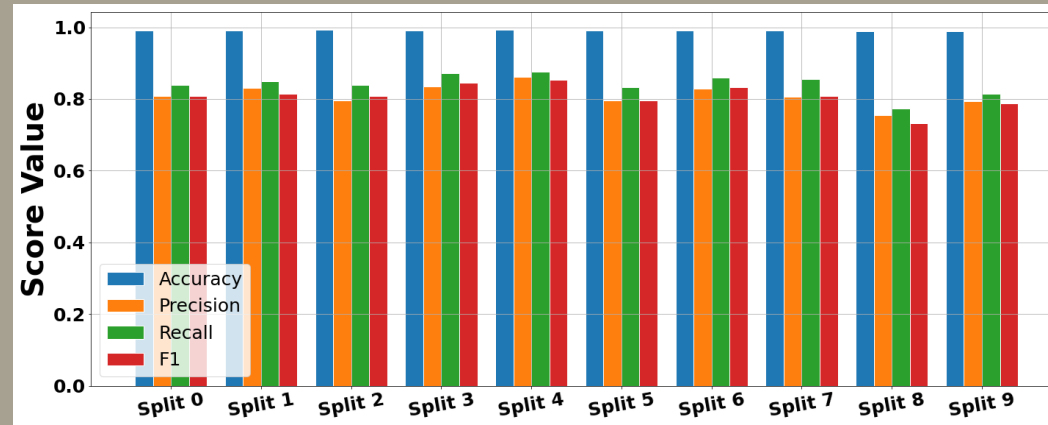


WE *TEST* THE MODEL ON THE  
REMAINING SAMPLES THAT THE  
MODEL HAS *NEVER INTERACTED*  
WITH

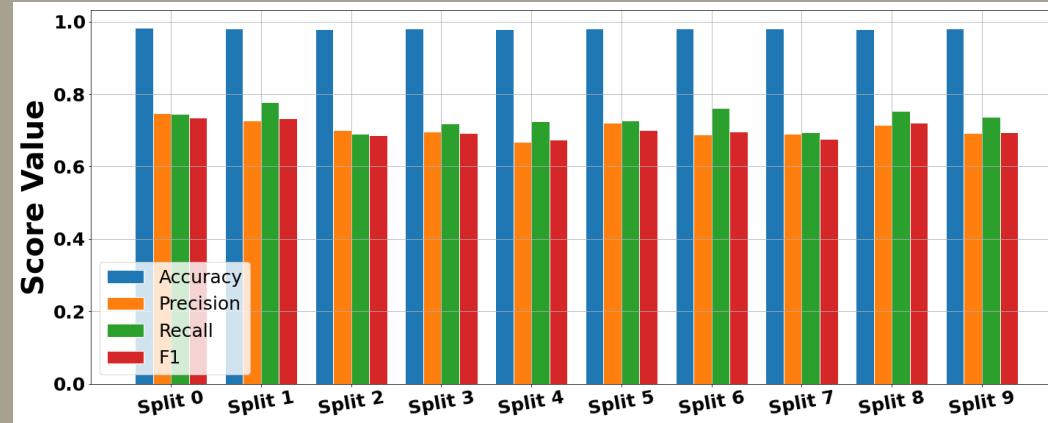
# X-Validation Performance

- Compare three variants
  - 1) CNN Architecture
  - 2) MLP Architecture
  - 3) Hybrid Architecture
- Hybrid network shows generally better and more consistent performance scores

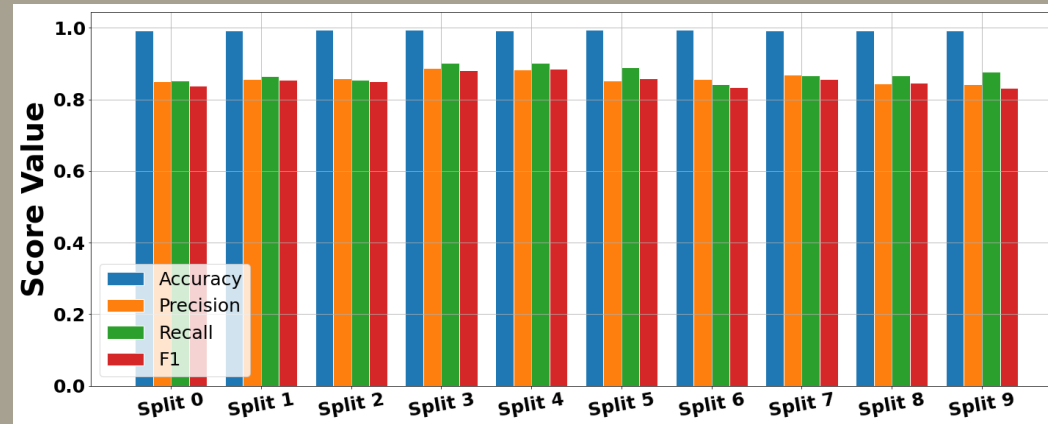
1)



2)

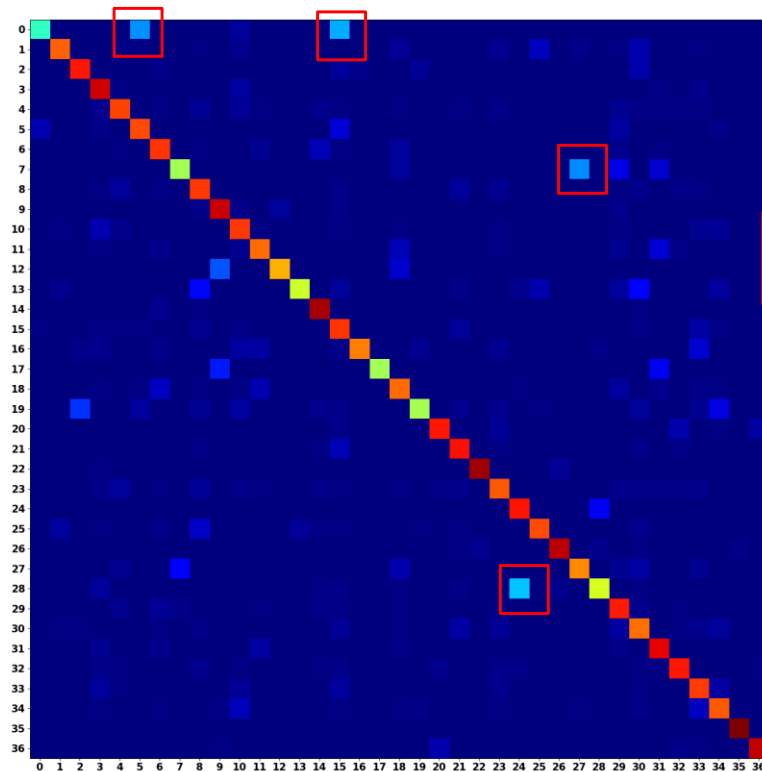


3)



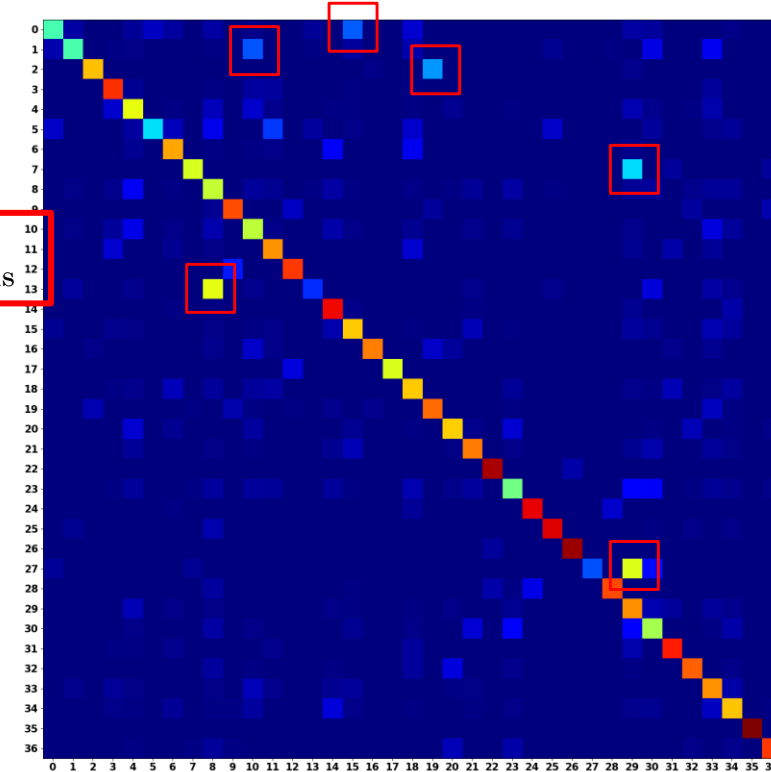
# Single NN Confusion Matrices

CNN Branch (Occurrence Weighted)



Strong main diagonal,  
reasonable performance

MLP Branch (Occurrence Weighted)



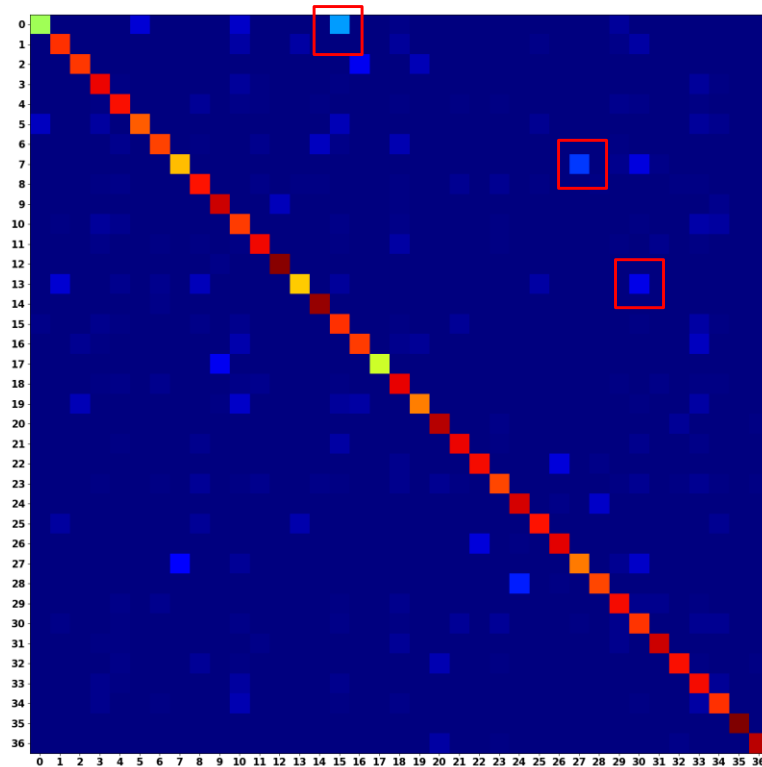
Frequent  
misclassifications

Related classes?  
Poor features choice?



# Hybrid NN Confusion Matrix

CNN + MLP (Occurrence Weighted)



- Hybrid model shows a stronger main diagonal
- Combines predictive power from both architectures

Less Frequent, Less  
Pronounced misclassifications

# Discussion

- Hybrid networks allows for *multimodal learning*
  - Transform inputs into compatible formats
  - Combine multiple inputs to form one prediction
- Opens the door to other related classification tasks
- Many uses outside digital audio recognition
  - Cardiograph + Biometric readings
  - Speech + text information
  - Video + Audio

- The *hybrid network* shows improved classification performance over either unimodal model
- Possibility to generalize to other related tasks

# Conclusions

# References

[1] Geron, Aurelien. Hands-on Machine Learning with Scikit-Learn and TensorFlow: Concepts, Tools, and Techniques to Build Intelligent Systems. O'Reilly, 2017.

[2] Goodfellow, Ian, et al. Deep Learning. MIT Press, 2017.

[3] Khan, M. Kashif Saeed, and Wasfi G. Al-Khatib. "Machine-Learning Based Classification of Speech and Music." Multimedia Systems, vol. 12, no. 1, 2006, pp. 55–67., doi:10.1007/s00530-006-0034-0.

[4] Li, Yingming, and Ming Yang. "A Survey of Multi-View Representation Learning." Journal of LateX Class Files, vol. 14, no. 8, Aug. 2015.

[5] Liu, Zhu, et al. "Audio Feature Extraction and Analysis for Scene Segmentation and Classification." Journal of VLSI Signal Processing, vol. 20, 1998, pp. 61–79.

[6] Ngiam, Jiquan, et al. "Multimodal Deep Learning." 2011.

[7] TensorFlow: Large-scale machine learning on heterogeneous systems, 2015. Software available from tensorflow.org.

[8] Virtanen, Tuomas, et al. Computational Analysis of Sound Scenes and Events. Springer, 2018.

[9] White, Harvey Elliott, and Donald H. White. Physics and Music: the Science of Musical Sound. Dover Publications, Inc., 2019.

Additional information, GitHub  
Repository, and Formal Write-Up is  
Available upon request

- Landon Buell – [lhb1007@wildcats.unh.edu](mailto:lhb1007@wildcats.unh.edu)
- Kevin Short – [kevin.short@unh.edu](mailto:kevin.short@unh.edu)

Thank you very  
much!

Questions?

# Appendix – Confusion Matrices

## Standard Confusion Matrix

- For  $k$  categories, is a  $k \times k$  matrix

$C[i, j]$  = Number of samples that belong to class  $i$  but were predicted to be in class  $j$

- A strong classifier has a dominant main-diagonal

## Occurrence Weighted Matrix

- Created same as standard matrix

- Divide each row by sum of the row

- Accounts for non-uniform number of samples in each class

# Appendix - Metrics

## Precision / Specificity Score

- Bound on  $[0,1]$  – higher is favorable

$$Precision = \frac{TP}{TP + FP}$$

- *“How many selected items are relevant to the problem?”*

## Recall / Sensitivity Score

- Bound on  $[0,1]$  – higher is favorable

$$Recall = \frac{TP}{TP + FN}$$

- *“How many relevant items to the problem have been selected?”*

# Appendix – Metrics (Cont.)

## Accuracy Score

- Ratio of correct predictions to total predictions

$$Accuracy = \frac{TP + FN}{TP + FP + FN + TP}$$

- Does not account for non-uniform number of samples in each class

## F1 Score

- Harmonic Mean of Precision and Recall scores

$$F1 = 2 \times \frac{Precision + Recall}{Precision \times Recall}$$

- Favors models with high precision *and* high recall



# Appendix – Activation Functions

- Activation functions are typically included as a “last” step in each layer function
- Can take many forms based on network or layer type
  - ReLU for classification
  - Sigmoid for regression
  - Softmax for normalized output
- Turns affine-transform into non-linear transform
- Model more complex solution spaces

# Appendix – Tensorflow / Keras

- Free Python library centered around Deep Learning
- High-Level API for constructing neural networks and computational graphs
- Developed and maintained by Google  
<https://www.tensorflow.org/>
- Offers tools for hardware acceleration on certain GPUs

# Appendix – Musical Instrument Samples

- University of Iowa Electronic Music Studios  
[University of Iowa Electronic Music Studios \(uiowa.edu\)](http://www.music.uiowa.edu/)
- Philharmonia Symphony Orchestra  
<https://philharmonia.co.uk/>
- Samples preprocessed
  - WAV file format
  - 44.1 kHz sample rate
  - Mono-channeled
  - 16 bit-depth
- Features extracted with Python program available on GitHub